



Frank Stale

Software Engineer & Builder

✉ frank.stale@example.com



+447402528867

[in/FrankStale](#)

[github](#)

Computer Science undergraduate at the University of Birmingham with a track record of delivering full-stack features in Spring Boot, React, and PostgreSQL across both internship and freelance settings. Comfortable across the stack — from data-pipeline work and ML prototyping in Python to containerised web applications in production. Looking to apply solid engineering fundamentals on a team that values clean code, mentorship, and tangible product impact.

Experience

Software Engineer at Company 1 — Apr 2021 – Present

Brightline Software — Birmingham, UK

- Owned a customer onboarding micro-service written in Node.js and Express, reducing average new-account setup time from twelve minutes to under three through batched API calls and caching.
- Built a React dashboard for the operations team to monitor sign-up funnels, integrating with the existing PostgreSQL warehouse and exposing trend charts, retention cohorts, and per-channel attribution.
- Partnered with QA to introduce contract testing across the public REST API, cutting regression incidents on releases by roughly a third over the following quarter.

Software Engineer at Company 2 — Oct 2020 – Present

Halcyon Health — Remote (UK)

- Contributed to a Spring Boot patient-records service handling roughly four million events per day, focusing on a refactor of the audit trail to use append-only event sourcing.
- Implemented a Hibernate-backed query layer with first-class soft-delete support, replacing scattered ad-hoc filters and removing a class of data-leakage bugs raised in clinical reviews.
- Wrote internal documentation and runbooks for the on-call rotation, which reduced the average time to acknowledge production alerts by approximately 20 minutes.

Lloyds Banking Group — Sep 2021 – Present

Technology & Data Insight Programme — Birmingham, UK

- Engineered predictive models to assess customer overdraft behaviour using transactional data, leveraging pandas and scikit-learn to streamline the data-processing pipeline.
- Achieved roughly 72% accuracy in identifying repeat overdraft customers via logistic regression, supporting downstream financial-risk decisioning.
- Created data visualisations with Matplotlib and presented findings to a panel of senior stakeholders, framing model behaviour in plain business terms.

Software Engineer at Company 4 — Aug 2018 – Present

Inkwell Studios — Bristol, UK

- Designed and shipped a Matplotlib-based reporting tool that ingested daily campaign data from a PostgreSQL store and produced PDF performance summaries for the marketing team.
- Optimised a slow data-cleaning pipeline by replacing row-level loops with vectorised pandas operations, taking a nightly job from forty minutes to under five.

- Presented findings on user retention to senior stakeholders, translating raw analytics into recommendations that informed two follow-up product experiments.

Projects

Educache — Aug 2024 – Present

Spring Boot, JPA, Hibernate, React, PostgreSQL, Docker, OpenAI API

Mobile Payment SDK — Jan 2021 – Sep 2024

Advanced News Classifier — Oct 2023 – Dec 2023

Java, Deeplearning4J, ND4J

Machine Learning Data Pipeline — Dec 2021 – May 2023

Volunteering — Nov 2022 – Feb 2024

Education

Bachelor of Science — 2023 – Feb 2020

Bristol, UK

Tools: Git — Nov 2018 – Mar 2020

Bristol, UK

Skills

C

Express

Haskell

Hibernate

Java

JavaScript

Matplotlib

Node.js

Pandas

PostgreSQL

Processing

Python

React

scikit-learn

Spring

SQL

TypeScript

Django