

Moa Lundh

+46 70 234 5678 | moa.lundh@gmail.com | <https://www.linkedin.com/in/moa-lundh>

EDUCATION

Royal Institute of Technology (KTH) — *M.S. in Computer Science, specialization in Distributed Systems* 2007 - 2009
Lund University — *B.S. in Computer Science and Mathematics* 2003 - 2007

EXPERIENCE

Spotify — *Principal Data Engineer* Jan 2021 - Present

- Architected and led adoption of real-time analytics platform processing 2.8B events/day across 400M users, reducing query latency by 64% and enabling 12 product teams to ship features 3x faster
- Established data quality framework using Monte Carlo Data and Great Expectations, reducing incidents by 82% and achieving 99.7% SLA compliance across 300+ critical pipelines
- Mentored 8 senior data engineers and 3 staff engineers on advanced streaming architecture and contributed 15 architectural decision records adopted org-wide

Uber — *Staff Data Engineer* Jun 2018 - Dec 2020

- Designed and shipped unified data lakehouse on Delta Lake handling 1.2PB of daily ingestion for 50+ teams, improving ETL reliability to 99.95% and reducing infrastructure costs by \$4.2M annually
- Built distributed feature store serving 800+ ML models with sub-100ms latency, enabling 40+ data scientists to reduce feature engineering time by 70% and accelerate model deployment cycles
- Led migration of 200+ legacy batch pipelines to Kafka + Spark Structured Streaming, eliminating 14-hour SLA window and enabling real-time decisioning for surge pricing and dispatch algorithms

Google Cloud — *Senior Data Engineer (Cloud Data Analytics)* Feb 2016 - May 2018

- Architected BigQuery optimization patterns adopted by 500+ enterprise customers, generating \$15M incremental ARR and published 3 technical papers on cost-efficient analytics at VLDB 2017
- Designed Dataflow-based ETL framework processing 200TB+ daily across financial services, healthcare, and media verticals with 99.98% accuracy and enabling 2,000+ analysts to run ad-hoc queries sub-second
- Established data governance and lineage platform serving 100+ GCP projects, reducing compliance audit time by 76% and enabling automated PII detection across 800+ datasets

Goldman Sachs — *Vice President, Data Engineering* Aug 2013 - Jan 2016

- Led platform team of 6 engineers to build real-time risk data pipeline ingesting 50TB daily from 15 asset classes, reducing risk calculation latency from 8 hours to 12 minutes and supporting \$2.3T trading volume
- Implemented horizontally-scalable data mesh for 30+ trading desks across equities, fixed income, and derivatives, enabling self-service BI access and reducing quant request queue by 84%
- Pioneered low-latency Kafka + Flink infrastructure for market data processing, handling 40M tick events/second with <50ms end-to-end latency for algorithmic trading strategies

Databricks — *Principal Solutions Architect, Data Platform* Sep 2011 - Jul 2013

- Architected Delta Lake adoption strategy for Fortune 100 companies spanning retail, insurance, and automotive; drove 8 successful implementations on 500TB+ datasets resulting in 52% improvement in data freshness
- Delivered 40+ technical workshops and conference presentations on Apache Spark optimization, data governance, and lakehouse best practices, reaching 8,000+ engineers and establishing thought leadership
- Authored open source framework for automated data lineage tracking across Spark jobs, now adopted by 30+ enterprises and featured as recommended tool in Apache Spark ecosystem

PROJECTS

DataLineage (Open Source) — *Python, Scala, Spark, Apache Atlas, GraphQL* 2012 - 2014

- Designed and open-sourced data lineage and impact analysis tool used by 2,000+ practitioners, achieving 800 GitHub stars and 95+ production deployments at scale
- Demonstrated automated end-to-end column-level lineage tracking and data cost attribution across 40M+ objects, enabling \$8M+ in infrastructure cost optimization for adopting enterprises

EXTRACURRICULAR

O'Reilly Media & Data Council — *Conference Speaker & Session Chair* 2015 - Present

- Keynoted Strata Data Conference 2018, 2019, 2021 on 'Scaling Real-Time Analytics to Billions of Events'
- Organized and chaired Data Infrastructure & Engineering Summit sessions, featuring talks from 50+ engineers reaching 3,000+ attendees

Swedish Data Engineering Society (SDES) — *Board Member & Program Advisor* 2017 - Present

- Established Nordic Data Engineering Working Group connecting 2,000+ practitioners across Nordics for knowledge sharing and best practice standardization
- Mentored 12 emerging data engineering leaders and helped 3 transition into Principal roles at top-tier companies

TECHNICAL SKILLS

Languages – Python, Scala, SQL, Java, Go, Bash

Technologies – Apache Spark, Kafka, Airflow, dbt, Snowflake, BigQuery, Delta Lake, Ray, Dataflow, Flink, Protobuf, Avro

Tools – Terraform, Docker, Kubernetes, Jenkins, Datadog, Monte Carlo Data, Great Expectations, Tableau, Looker