

# Neuralflow AI

*Democratizing machine learning for enterprise applications*

## INTERN MACHINE LEARNING ENGINEER

---

San Francisco, CA (Hybrid) · Internship · \$22-26 per hour

We are seeking a talented Machine Learning Engineer Intern to join our ML Infrastructure team. In this role, you will gain hands-on experience building and optimizing machine learning models, working with production ML pipelines, and collaborating with experienced engineers to solve real-world ML challenges. You'll contribute to projects spanning model training, deployment, and monitoring while learning best practices in machine learning engineering.

## ABOUT US

---

Neuralflow AI is a fast-growing machine learning platform company based in San Francisco that helps enterprises build, train, and deploy production-grade ML models at scale. We provide end-to-end ML infrastructure and tooling that simplifies the entire ML lifecycle, from data preparation to model monitoring. Our mission is to make advanced machine learning accessible to organizations of all sizes.

## WHAT YOU'LL DO

---

- Assist in developing and training machine learning models using PyTorch and TensorFlow on real-world datasets
- Help implement feature engineering pipelines using Featuretools and collaborate on data processing tasks
- Participate in model experimentation and tracking using MLflow to log metrics, parameters, and model artifacts
- Support the deployment and testing of ML models on AWS SageMaker, including managing model versions and inference endpoints
- Write clean, well-documented Python code and participate in code reviews with the team
- Monitor model performance metrics and alert team members to potential data drift or model degradation
- Contribute to documentation and create sample notebooks demonstrating ML workflows for internal use
- Attend team meetings, learning sessions, and present findings on model experiments and improvements

## REQUIRED

---

- Currently pursuing or recently completed a Bachelor's degree in Computer Science, Data Science, Statistics, Mathematics, or related field
- Strong programming skills in Python with solid understanding of object-oriented programming principles

- Foundational knowledge of machine learning concepts including supervised learning, classification, and regression
- Familiarity with at least one ML framework (TensorFlow, PyTorch, or scikit-learn) through coursework or personal projects

## NICE TO HAVE

---

- Prior internship or project experience with machine learning or data science applications
- Experience with cloud platforms (AWS, Google Cloud Platform, or Azure) or familiarity with containerization (Docker)
- Knowledge of ML operations (MLOps) best practices, model versioning, or experiment tracking tools

## TECH STACK

---

Python

PyTorch

TensorFlow

MLflow

AWS SageMaker

Featuretools

Pandas & NumPy

Docker

Git/GitHub

Jupyter Notebooks

## BENEFITS

---

- Competitive hourly internship compensation with potential for return offer
- Flexible work schedule accommodating academic calendar and course load
- Comprehensive onboarding and mentorship from senior ML engineers
- Access to cutting-edge ML tools, GPUs, and cloud computing resources

## HOW TO APPLY

---

Please submit your resume, a cover letter describing your interest in machine learning, and a link to your GitHub profile or portfolio (if available) to [careers@neuralflow.ai](mailto:careers@neuralflow.ai) with the subject line 'ML Engineer Intern Application'. Applications are reviewed on a rolling basis. We encourage students and recent graduates to apply!