

DataFlow Innovations

Building the data infrastructure that powers real-time insights at scale

LOCATION

San Francisco, CA (Remote-Friendly)

TYPE

Full-Time

COMPENSATION

\$240,000 - \$320,000 USD

TECH STACK

- Apache Spark
- Airflow
- dbt
- Snowflake
- Python
- SQL
- Kafka
- BigQuery
- Data modeling (dimensional modeling, data vault)
- Terraform/Infrastructure-as-Code

Staff Data Engineer

About DataFlow Innovations: DataFlow Innovations is a fast-growing data infrastructure company founded in 2018, specializing in building enterprise-grade data platforms and analytics solutions for mid-market to Fortune 500 companies. We help organizations ingest, transform, and operationalize massive volumes of data across hybrid cloud environments. Our platform serves over 200 customers across finance, healthcare, retail, and technology sectors.

We're seeking an experienced Staff Data Engineer with 8+ years of proven technical leadership to join our Data Platform team. In this role, you'll architect and lead the design of large-scale data pipelines, mentor a team of senior engineers, and establish technical standards for data infrastructure across the organization. You'll work on mission-critical data systems supporting thousands of concurrent users while guiding architectural decisions that impact company-wide data strategy.

RESPONSIBILITIES

- Architect and design enterprise-scale data platforms supporting petabyte-level data ingestion, transformation, and analytics workloads using Apache Spark, Airflow, and cloud data warehouses
- Lead technical design reviews and provide architectural guidance for complex data pipeline projects across multiple teams, establishing best practices for data modeling, lineage tracking, and quality assurance
- Mentor and develop a team of 4-6 senior data engineers, conducting code reviews, fostering technical growth, and building a culture of excellence and knowledge-sharing
- Establish and enforce data engineering standards, including SQL best practices, dbt conventions, infrastructure-as-code patterns, and CI/CD workflows for data products
- Drive the adoption of new technologies and frameworks (e.g., Kafka for real-time streaming, advanced dbt patterns) by conducting POCs, building reference implementations, and training engineering teams
- Collaborate with cross-functional stakeholders including data scientists, analytics engineers, and product teams to translate business requirements into scalable technical solutions

- Improve data platform reliability and performance by identifying bottlenecks in Snowflake/BigQuery queries, optimizing Airflow DAGs, and implementing monitoring and alerting strategies
- Lead strategic initiatives for data platform modernization, including cloud migration planning, disaster recovery design, and cost optimization across data infrastructure
- Contribute to open-source projects and present technical work at internal tech talks and industry conferences

REQUIREMENTS

- 8+ years of hands-on data engineering experience with proven track record designing and shipping production data systems at scale
- Expert-level proficiency in Python and SQL, with ability to write performant queries and optimize complex transformations
- 5+ years of production experience with Apache Spark and workflow orchestration tools (Airflow, Prefect, or Dagster), with deep understanding of distributed computing concepts
- 3+ years of experience building and optimizing cloud data warehouse solutions (Snowflake, BigQuery, Redshift) including schema design, partitioning strategies, and cost optimization
- Demonstrated technical leadership: mentoring engineers, driving architectural decisions, and influencing technical direction across teams
- Strong systems thinking and experience building data lineage, quality frameworks, and metadata management systems

BONUS POINTS

- Experience with dbt (data build tool) and modern ELT/dbt best practices for data transformation at scale
- Real-time data pipeline experience using Kafka, Spark Streaming, or similar event streaming platforms
- Hands-on experience with Iceberg, Delta Lake, or other open table formats for data lakehouse architectures
- Knowledge of data governance, data quality frameworks (Great Expectations, dbt tests), and regulatory compliance (GDPR, HIPAA)
- Track record of successful hiring, team building, and organizational scaling

BENEFITS

- Competitive salary and equity package with annual refreshing grants
- Comprehensive health coverage (medical, dental, vision) with 90% employer premium coverage
- 401(k) match up to 6% with immediate vesting
- Unlimited PTO with minimum 15 days encouraged per year
- Remote-first work environment with flexible office access in SF, NYC, and Austin
- Annual professional development budget (\$5,000) for conferences, certifications, and courses
- Parental leave (16 weeks primary caregiver, 12 weeks secondary)

APPLY

Please submit your resume, a portfolio or GitHub profile showcasing your data engineering work, and a brief cover letter explaining your interest in data platform architecture to careers@dataflow-innovations.com. In your application, highlight your experience with large-scale data systems and any leadership roles you've held. We review applications on a rolling basis and aim to respond within 1 week.