

Manager, MLOps Engineer

VectorFlow AI — Democratizing machine learning operations at enterprise scale

Location	San Francisco, CA (Remote-friendly)
Employment Type	Full-time
Compensation	\$210,000 - \$270,000 + equity

The Opportunity

Join VectorFlow AI as Manager, MLOps Engineer and lead a team of 5-7 engineers responsible for building and scaling our core MLOps infrastructure. In this role, you'll oversee the architecture and deployment of ML workflows supporting hundreds of production models, manage a high-performing team, mentor engineers on MLOps best practices, and drive strategic initiatives around model observability and platform reliability. You'll work closely with product, data science, and infrastructure teams to ensure our platform delivers exceptional performance and reliability to enterprise customers.

About VectorFlow AI

VectorFlow AI is a Series B-funded MLOps platform company dedicated to helping enterprises streamline their machine learning infrastructure, model deployment, and production monitoring. We serve 200+ companies across financial services, healthcare, and e-commerce, enabling data teams to go from experimentation to production in days, not months. Our platform integrates seamlessly with existing data stacks and provides end-to-end visibility into ML systems.

Day-to-Day

- Lead and manage a team of 5-7 MLOps engineers, including hiring, performance reviews, career development, and fostering a culture of technical excellence and collaboration
- Oversee the design and implementation of scalable ML infrastructure using Kubernetes, Docker, and cloud-native technologies to support 500+ production models across customer deployments
- Drive architectural decisions for workflow orchestration using Apache Airflow and Kubeflow, ensuring reliability, scalability, and ease of use for data teams
- Establish and enforce MLOps best practices, including CI/CD pipelines for model deployment, automated testing frameworks, and infrastructure-as-code standards using GitOps principles
- Build and maintain comprehensive model monitoring and governance solutions using MLflow, Prometheus, and custom observability tools to detect model drift, performance degradation, and data quality issues
- Partner with the Data Engineering team to integrate feature stores and manage feature pipelines that serve both training and inference workflows at scale
- Identify and execute on strategic technical roadmap items; propose and lead engineering initiatives that improve platform reliability, reduce operational toil, and enhance customer experience
- Conduct regular 1-on-1s with direct reports, provide technical mentorship, and support team members in advancing their technical skills and career growth

- Represent the MLOps team in cross-functional discussions with product, sales, and customer success to translate business requirements into technical roadmap priorities

You Have

- 6+ years of experience in MLOps, platform engineering, or machine learning infrastructure, with at least 2+ years in a management or tech lead role
- Demonstrated expertise building and operating production ML systems at scale, including model training pipelines, serving infrastructure, and monitoring systems
- Deep hands-on experience with at least two of: Apache Airflow, Kubeflow, or similar orchestration platforms; and proficiency with Docker and Kubernetes for containerized ML workloads
- Strong software engineering practices including version control, testing, CI/CD automation, and infrastructure-as-code (Terraform, CloudFormation, or similar)
- Proven track record managing and mentoring engineering teams, with experience in hiring, performance management, and creating psychological safety within teams
- Solid understanding of ML lifecycle concepts including model evaluation, hyperparameter optimization, model registry/versioning (e.g., MLflow), and model governance
- Experience with monitoring, observability, and alerting for ML systems; familiarity with tools like Prometheus, Grafana, or equivalent platforms

Bonus

- Experience building or maintaining a feature store solution or working closely with feature engineering platforms (e.g., Tecton, Feast, or internal solutions)
- Exposure to model monitoring frameworks and tools for detecting model drift, data drift, and prediction quality degradation
- Familiarity with cloud platforms (AWS, GCP, or Azure) and their managed ML services (SageMaker, Vertex AI, etc.)
- Background in data engineering or data infrastructure that complements MLOps knowledge
- Experience with open-source MLOps tools and active contribution to or leadership in the MLOps community

Our Stack



What We Offer

- Competitive salary and equity package; 401(k) matching (up to 4%)
- Comprehensive health coverage including medical, dental, vision, and mental health support
- Unlimited PTO with a minimum of 15 days encouraged; 10 company holidays
- Professional development budget (\$3,000/year for courses, conferences, and certifications)
- Remote-friendly culture with co-working stipend and option for office space in SF, NYC, or Seattle
- Parental leave (16 weeks primary caregiver, 12 weeks secondary caregiver) and fertility benefits

- Home office setup allowance and annual equipment refresh budget
- Team off-sites and regular social events; commuter benefits and gym membership reimbursement

Apply

Apply directly via our careers page at vectorflow.ai/careers or email your resume and a brief cover letter explaining your interest in MLOps leadership to careers@vectorflow.ai. We review applications on a rolling basis and aim to respond to all candidates within 2 weeks. Interviews will include a technical discussion, system design conversation, and team meeting with current MLOps leaders.